

Snowflake External Tables with S3 Integration

S3 Integration + Stages +Ext. table

Set-up:

Step 1: Configure access permissions for the S3 bucket:

Snowflake requires the following permissions on an S3 bucket and folder to be able to access files in the folder (and sub-folders):

- s3:GetBucketLocation
- s3:GetObject
- s3:GetObjectVersion
- s3>ListBucket

Creating an IAM policy:

- Log into the AWS Management Console.
- From the home dashboard, search for and select IAM.
- From the left-hand navigation pane, select Account settings.
- Under Security Token Service (STS) in the Endpoints list, find the Snowflake region where your account is located. If the STS status is inactive, move the toggle to Active.
- From the left-hand navigation pane, select Policies.
- Select Create Policy.
- For Policy editor, select JSON.

- Copy and paste the text into the policy editor (replace bucket and prefix with your actual bucket name and folder path prefix, Setting the "s3:prefix": condition to either ["*"] or ["<path>/*"] grants access to all prefixes in the specified bucket or path in the bucket) :

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:PutObject",
        "s3:GetObject",
        "s3:GetObjectVersion",
        "s3:DeleteObject",
        "s3:DeleteObjectVersion"
      ],
      "Resource": "arn:aws:s3:::<bucket>/<prefix>/*"
    },
    {
      "Effect": "Allow",
      "Action": [
        "s3>ListBucket",
        "s3:GetBucketLocation"
      ],
      "Resource": "arn:aws:s3:::<bucket>",
      "Condition": {
        "StringLike": {
          "s3:prefix": [
            "<prefix>/*"
          ]
        }
      }
    }
  ]
}
```

- Select Next.
- Enter a Policy name (for example, snowflake_access) and an optional Description.
- Select Create policy.

Step 2: Create the IAM role in AWS

- From the left-hand navigation pane in the Identity and Access Management (IAM) Dashboard, select Roles.
- Select Create role.

- Select AWS account as the trusted entity type.
- Select Another AWS account

The screenshot shows the AWS IAM 'Create role' wizard, Step 1: Select trusted entity. The 'AWS account' option is selected, highlighted with a blue border. Other options include 'AWS service', 'Web identity', 'SAML 2.0 federation', and 'Custom trust policy'. Below this, there's a section for 'An AWS account' where 'Another AWS account' is selected. An account ID '123456789012' is entered in the 'Identifier of the account that can use this role' field. Under 'Options', 'Require external ID' is checked, and an external ID '0000' is entered. A note states: 'Important: The console does not support using an external ID with the Switch Role feature. If you select this option, entities in the trusted account must use the API, CLI, or a custom federation proxy to make cross-account iam:AssumeRole calls.' At the bottom, there are 'Cancel' and 'Next' buttons.

- In the Account ID field, enter your own AWS account ID temporarily. Later, you modify the trust relationship and grant access to Snowflake.

- Select the Require external ID option. An external ID is used to grant access to your AWS resources (such as S3 buckets) to a third party like Snowflake.
- Enter a placeholder ID such as 0000. In a later step, you will modify the trust relationship for your IAM role and specify the external ID for your storage integration.
- Select Next.
- Select the policy you created in Step 1: Configure access permissions for the S3 bucket (in this topic).
- Select Next.
- Enter a name and description for the role, then select Create role.
- You have now created an IAM policy for a bucket, created an IAM role, and attached the policy to the role.
- On the role summary page, locate and record the Role ARN value. In the next step, you will create a Snowflake integration that references this role.

Step 3: Create a cloud storage integration in Snowflake

Create a storage integration using the CREATE STORAGE INTEGRATION command:

```
CREATE OR REPLACE STORAGE INTEGRATION <give integration a name>
  TYPE = EXTERNAL_STAGE
  STORAGE_PROVIDER = 'S3'
  ENABLED = TRUE
  STORAGE_AWS_ROLE_ARN = '<iarn role arn>'
  STORAGE_ALLOWED_LOCATIONS = ('*');
```

Step 4: Retrieve the AWS IAM user for your Snowflake account

- To retrieve the ARN for the IAM user that was created automatically for your Snowflake account, use the DESCRIBE INTEGRATION:
“*DESC INTEGRATION <integration_name>*”;
- Record the values for the following properties:

“*STORAGE_AWS_IAM_USER_ARN*”:

The AWS IAM user created for your Snowflake account; for example, `arn:aws:iam::123456789001:user/abc1-b-self1234`. Snowflake provisions a single IAM user for your entire Snowflake account. All S3 storage integrations in your account use that IAM user.

“*STORAGE_AWS_EXTERNAL_ID*”:

The external ID that Snowflake uses to establish a trust relationship with AWS. If you didn’t specify an external ID (STORAGE_AWS_EXTERNAL_ID) when you created the storage integration, Snowflake generates an ID for you to use.

Step 5: Grant the IAM user permissions to access bucket objects

- Select the role you created in Step 2: Create the IAM role in AWS (in this topic).
- Select the Trust relationships tab.
- Select Edit trust policy.
- Modify the policy document with the DESC STORAGE INTEGRATION output values you recorded in Step 4: Retrieve the AWS IAM user for your Snowflake account (in this topic):
- Policy document for IAM role

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "",
      "Effect": "Allow",
      "Principal": {
        "AWS": "<snowflake_user_arn>"}
```

```

        },
        "Action": "sts:AssumeRole",
        "Condition": {
            "StringEquals": {
                "sts:ExternalId": "<snowflake_external_id>"
            }
        }
    }
}

```

- snowflake_user_arn is the STORAGE_AWS_IAM_USER_ARN value you recorded.
- snowflake_external_id is the STORAGE_AWS_EXTERNAL_ID value you recorded.

Step 6: Create an external stage

```

CREATE OR REPLACE STAGE s3StageForSOD
    STORAGE_INTEGRATION = s3IntegrationForSOD
    URL = 's3://niinebucket/folder_for_sod/'
    FILE_FORMAT = (TYPE = 'CSV' FIELD_DELIMITER = '\t' SKIP_HEADER = 1);

```

Replace 'URL' with desired s3 folder, 'FILE_FORMAT' can also be changed according to CSV type **but the file should be of CSV only** as Snowflake allows only CSV external tables.

- Example file format:

```

create or replace file format CsvFormatForSodCSV
    type = 'csv'
    compression = 'auto'
    field_delimiter = ','
    record_delimiter = '\n'
    skip_header = 1
    field Optionally_enclosed_by = '\042'
    null_if = ('\\N');

```

Step 7: Create an external table

Below is an example format for external table that casts all columns as VARCHAR for avoiding value datatype errors:

```
CREATE OR REPLACE EXTERNAL TABLE S3_DATA_NIINE.PUBLIC.NewSodExtTable (
Order_No VARCHAR(16777216) AS (value:c1::VARCHAR),
ESM_Erp_Id VARCHAR(16777216) AS (value:c2::VARCHAR),
User_Designation VARCHAR(16777216) AS (value:c3::VARCHAR),
User_Position_Level VARCHAR(16777216) AS (value:c4::VARCHAR),
User VARCHAR(16777216) AS (value:c5::VARCHAR),
SuperStockist VARCHAR(16777216) AS (value:c6::VARCHAR),
SuperStockist_Erpld VARCHAR(16777216) AS (value:c7::VARCHAR),
Distributor_Erp_Id VARCHAR(16777216) AS (value:c8::VARCHAR),
Beat_Erp_Id VARCHAR(16777216) AS (value:c9::VARCHAR),
Beat VARCHAR(16777216) AS (value:c10::VARCHAR),
Outlet_Erp_Id VARCHAR(16777216) AS (value:c11::VARCHAR),
Outlet_Name VARCHAR(16777216) AS (value:c12::VARCHAR),
Order_Date VARCHAR(16777216) AS (value:c13::VARCHAR),
Price VARCHAR(16777216) AS (value:c14::VARCHAR),
Sale_Value VARCHAR(16777216) AS (value:c15::VARCHAR),
Call_Start_Time VARCHAR(16777216) AS (value:c16::VARCHAR),
Call_End_Time VARCHAR(16777216) AS (value:c17::VARCHAR),
Product_Erpld VARCHAR(16777216) AS (value:c18::VARCHAR),
Product VARCHAR(16777216) AS (value:c19::VARCHAR),
Qty_StdUnit VARCHAR(16777216) AS (value:c20::VARCHAR),
Gross_Value VARCHAR(16777216) AS (value:c21::VARCHAR)
)
WITH LOCATION = @s3StageForSOD
FILE_FORMAT = 'CsvFormatForSodCSV'
PATTERN = '.*Secondary_Order_Dump_New_india.*\\.csv$'
AUTO_REFRESH = TRUE;
```

NOTE: Column names should not have spaces in them, if they are present then use underscores ('_') or enclose the column name in double quotes, for eg: “*Gross Value*” or “*Gross_Value*”